

Dual Privacy Architecture: Information-Theoretic Bounds on Agent Reconstruction

Mathematical Framework for Swordsman-Mage Separation

privacymage

November 25, 2025

Version 3.2

Abstract

We introduce the Swordsman and Mage as fundamental privacy primitives for dual-agent architectures, establishing rigorous information-theoretic bounds when conditional independence $(Y_S \perp Y_M) \mid X$ is enforced between these agents' observations. The Swordsman (S) controls privacy boundaries through selective measurement, while the Mage (M) projects delegated agency using only S-authorized observations.

Proven Results: We prove that this separation enables an additive bound on mutual information: $I(X; Y_S, Y_M) \leq I(X; Y_S) + I(X; Y_M)$. Combined with budget constraints $C_S + C_M < H(X)$, this establishes a reconstruction ceiling $R_{\max} < 1$ that no adversary can exceed regardless of computational resources. Via Fano's inequality, we establish a fundamental error floor: $P_e \geq 1 - (I(X; Y) + 1)/H(X)$, guaranteeing minimum reconstruction error when $R_{\max} < 1$. We further prove graceful degradation under approximate separation.

Implementation Framework: We provide practical budget estimation methods, isolation verification protocols, and side-channel resistance models based on covert channel analysis. We integrate zero-knowledge proof systems for cryptographic enforcement of separation and budget compliance, providing concrete constructions using Groth16, PLONK, and Nova protocols.

Theoretical Predictions: We present theoretical conjectures about potential optimal allocation patterns, including a golden ratio hypothesis ($\phi \approx 1.618$) and tetrahedral emergence properties. These remain unproven mathematical conjectures requiring both formal proof and empirical validation. We outline specific research directions and invite collaboration on implementation and validation.

Nature of This Work

What Is Proven: The core information-theoretic results (additive bounds under separation, reconstruction ceilings, error floors) are rigorously proven using established information theory.

What Is Theoretical: The golden ratio optimization hypothesis and tetrahedral emergence predictions are unproven mathematical conjectures. They represent interesting theoretical possibilities but have not been formally derived from first principles.

What Is Missing: No implementations exist. No empirical data has been collected. No observations have been made. This is purely theoretical and mathematical work at present.

What We Seek: Collaboration from theorists to prove or disprove the conjectures, and from practitioners to build implementations and collect empirical data.

1 Introduction

1.1 Motivation

The deployment of autonomous AI agents acting on behalf of humans creates a fundamental tension: agents require information about their principals to act effectively (delegation), yet this same information enables reconstruction of sensitive behavioral patterns (privacy loss). Traditional single-agent architectures cannot resolve this tension—the same system handling both functions creates an inherent conflict of interest.

We propose the Swordsman and Mage as dual privacy primitives that resolve this tension through architectural separation:

- **The Swordsman (S):** A privacy-enforcement primitive that controls information disclosure through selective measurement, acting as the guardian of boundaries
- **The Mage (M):** A delegation primitive that projects agency into the world using only Swordsman-authorized observations, acting as the executor of capabilities

These primitives are not mere metaphors but formal architectural components with precise information-theoretic properties. The key insight: enforcing conditional independence between the Swordsman and Mage observations creates provable reconstruction bounds.

1.2 Contributions

Proven Results:

- **Separation Lemma (Theorem 3.1):** Under $(Y_S \perp Y_M) \mid X$, mutual information becomes additive
- **Reconstruction Ceiling (Corollary 3.2):** With $C_S + C_M < H(X)$, reconstruction efficiency $R_{\max} < 1$
- **Error Floor (Theorem 3.3):** Fano’s inequality establishes minimum error $P_e \geq 1 - (I(X; Y) + 1)/H(X)$
- **Robustness Analysis (Theorem 3.4):** ε -approximate separation degrades bounds gracefully

Implementation Framework:

- Practical budget estimation and monitoring methods
- Isolation verification and enforcement protocols
- Side-channel resistance model with concrete parameters
- Zero-knowledge proof constructions for cryptographic separation enforcement
- ZKP-based budget compliance verification protocols
- Concrete specifications for Groth16, PLONK, and Nova-based implementations

Theoretical Predictions (unproven conjectures):

- Golden ratio convergence hypothesis in optimal allocations
- Tetrahedral emergence hypothesis for system architecture

1.3 Related Work

This work differs from existing privacy frameworks:

- **Differential Privacy** [1]: Adds calibrated noise; we enforce structural separation
- **Secure Multi-Party Computation** [2]: Distributes computation; we distribute observation rights
- **Information Flow Control** [3]: Tracks taint; we bound reconstruction
- **Covert Channel Analysis** [8]: Models side-channels; we apply to privacy architectures
- **Zero-Knowledge Proofs** [23, 24]: Verifiable computation; we apply to privacy budget enforcement

2 Model and Preliminaries

2.1 Basic Framework

Let X be a secret over finite alphabet $\mathcal{X}$ with $H(X) > 0$. Two agents produce observations:

$$Y_S = E_S(X, N_S) \tag{1}$$

$$Y_M = E_M(X, N_M) \tag{2}$$

where N_S, N_M are independent local randomness sources.

2.2 The Swordsman and Mage Primitives

Definition 2.1 (Swordsman Primitive). The Swordsman S is a privacy-enforcement agent characterized by:

- Measurement function E_S that implements selective disclosure
- Information budget C_S controlling maximum leakage: $I(X; Y_S) \leq C_S$
- Primary objective: minimize reconstruction while enabling necessary delegation

Definition 2.2 (Mage Primitive). The Mage M is a delegation agent characterized by:

- Projection function E_M operating on S-authorized information
- Information budget C_M for capability execution: $I(X; Y_M) \leq C_M$
- Primary objective: maximize utility under privacy constraints

The critical architectural requirement: $(Y_S \perp Y_M) \mid X$ (conditional independence).

2.3 Formal Definitions

Definition 2.3 (Separation Condition). The architecture enforces $(Y_S \perp Y_M) \mid X$.

Definition 2.4 (Information Budgets). $I(X; Y_S) \leq C_S, I(X; Y_M) \leq C_M$.

Definition 2.5 (Reconstruction Efficiency). $R \triangleq I(X; Y)/H(X) \in [0, 1]$.

2.4 Threat Model

Assumptions:

- Passive adversary observing (Y_S, Y_M)
- Separation enforced through architectural boundaries
- Known distributions $P(X)$, encoding functions E_S, E_M

Explicitly Out of Scope:

- Active attacks modifying agent behavior
- Side-channels on separation mechanism itself
- Temporal correlation across sessions

3 Proven Information-Theoretic Results

3.1 The Separation Lemma

Theorem 3.1 (Additive Bound Under Separation). If $(Y_S \perp Y_M) \mid X$ holds, then:

$$I(X; Y_S, Y_M) \leq I(X; Y_S) + I(X; Y_M)$$

Note: Equality holds if and only if Y_S and Y_M are additionally marginally independent. Conditional independence $(Y_S \perp Y_M) \mid X$ alone permits dependence through the shared cause X , yielding the inequality. The inequality suffices for all downstream guarantees.

Proof:

By the chain rule for mutual information:

$$I(X; Y_S, Y_M) = I(X; Y_S) + I(X; Y_M \mid Y_S)$$

Under conditional independence $(Y_S \perp Y_M) \mid X$, we have $I(Y_M; Y_S \mid X) = 0$, which implies:

$$H(Y_M \mid X, Y_S) = H(Y_M \mid X)$$

Therefore:

$$I(X; Y_M \mid Y_S) = H(Y_M \mid Y_S) - H(Y_M \mid X, Y_S) \tag{3}$$

$$= H(Y_M \mid Y_S) - H(Y_M \mid X) \tag{4}$$

$$\leq H(Y_M) - H(Y_M \mid X) \tag{5}$$

$$= I(X; Y_M) \tag{6}$$

This completes the proof. $\square$

Corollary 3.1.1 (Reconstruction Ceiling). If $C_S + C_M < H(X)$, then $R_{\max} = (C_S + C_M)/H(X) < 1$.

Critical Clarification: Separation alone is insufficient. Consider X binary with $Y_S = X$ and Y_M independent noise. Then $(Y_S \perp Y_M) \mid X$ holds but $R = 1$. The ceiling requires **BOTH** separation (for additivity) **AND** budget constraints (for the bound).

3.2 Error Lower Bound

Theorem 3.2 (Fano-Based Error Floor). For any estimator $\hat{X}(Y)$ and finite alphabet $\mathcal{X}$:

$$P_e \triangleq \Pr[\hat{X}(Y) \neq X] \geq \frac{H(X | Y) - 1}{\log(|\mathcal{X}| - 1)}$$

Which can be rewritten using $I(X; Y) = H(X) - H(X | Y)$ as:

$$P_e \geq \frac{H(X) - I(X; Y) - 1}{\log(|\mathcal{X}| - 1)}$$

For large alphabets where $\log(|\mathcal{X}| - 1) \approx H(X)$:

$$P_e \gtrsim 1 - \frac{I(X; Y) + 1}{H(X)} = 1 - R - \frac{1}{H(X)}$$

Proof:

Apply Fano's inequality in its standard form. For any estimator $\hat{X}(Y)$, the conditional entropy $H(X | Y)$ bounds the error probability:

$$H(X | Y) \leq h(P_e) + P_e \cdot \log(|\mathcal{X}| - 1)$$

where $h(\cdot)$ is binary entropy. Since $h(P_e) \leq 1$, we have:

$$H(X | Y) \leq 1 + P_e \cdot \log(|\mathcal{X}| - 1)$$

Rearranging:

$$P_e \geq \frac{H(X | Y) - 1}{\log(|\mathcal{X}| - 1)}$$

Using $I(X; Y) = H(X) - H(X | Y)$ and $R = I(X; Y)/H(X)$:

$$P_e \geq \frac{H(X) - I(X; Y) - 1}{\log(|\mathcal{X}| - 1)} \tag{7}$$

$$\geq 1 - \frac{I(X; Y) + 1}{H(X)} \quad [\text{for large alphabets}] \tag{8}$$

$$= 1 - R - \frac{1}{H(X)} \tag{9}$$

Therefore, when $R_{\max} < 1$, there exists a fundamental error floor that no estimator can overcome. $\square$

Interpretation: This establishes that when reconstruction efficiency $R_{\max} < 1$ (due to our budget constraints), any adversary must have error probability at least $P_e \geq 1 - R_{\max} - O(1/H(X))$. For example, if $R_{\max} = 0.7$, then $P_e \geq 0.3 - O(1/H(X)) \approx 0.3$ for large entropy.

3.3 Robustness to Approximate Separation

Theorem 3.3 (ε -Approximate Separation). If $I(Y_S; Y_M | X) \leq \varepsilon$ (approximate separation), then:

$$I(X; Y_S, Y_M) \leq I(X; Y_S) + I(X; Y_M) + \varepsilon$$

Proof:

Following the chain rule decomposition:

$$I(X; Y_M | Y_S) = H(Y_M | Y_S) - H(Y_M | X, Y_S)$$

With approximate independence:

$$H(Y_M | X, Y_S) \geq H(Y_M | X) - I(Y_S; Y_M | X) \geq H(Y_M | X) - \varepsilon$$

Therefore:

$$I(X; Y_M | Y_S) \leq H(Y_M) - H(Y_M | X) + \varepsilon = I(X; Y_M) + \varepsilon$$

Combining with $I(X; Y_S, Y_M) = I(X; Y_S) + I(X; Y_M | Y_S)$ completes the proof. $\square$

This shows graceful degradation: small violations of separation cause proportionally small increases in leakage.

4 Side-Channel Analysis and Robustness

4.1 Connection to Covert Channel Theory

Our analysis builds on established covert channel capacity results. For d side-channel observations, we model reconstruction growth as:

$$R(d) = R_{\max} \cdot \frac{\ln(1 + d/d_0)}{\ln(1 + d_{\max}/d_0)}$$

Model Assumption: This logarithmic functional form is adopted as a modeling assumption based on theoretical parallels, not derived from first principles. Empirical validation is required for specific deployment contexts.

This mirrors:

- **Shannon capacity** of noisy channels [6]
- **Timing channel capacity** under sampling [10]
- **Power analysis leakage** with noise [11]

4.2 Justification for Logarithmic Growth

The logarithmic form emerges from:

- **Information-theoretic argument:** Diminishing information per sample due to:
 - Temporal correlation in behavioral patterns
 - Finite substrate entropy $H(X)$
 - Measurement quantization effects
- **Theoretical foundation:** Previous theoretical work shows logarithmic scaling in:
 - Cache timing attacks [12]
 - Power analysis [13]
 - Traffic correlation [14]

4.3 Validation Methodology

To verify theoretical guarantees in practice:

- **Simulation Framework:**
 - Generate test distributions with known $H(X)$
 - Implement S and M with controlled coupling
 - Measure actual $I(X; Y_S, Y_M)$ vs theoretical bounds
- **Violation Detection Protocol:**
 - Monitor $I(Y_S; Y_M | X)$ in real-time
 - Flag when coupling exceeds threshold ε
 - Trigger corrective isolation measures

4.4 Parameter Estimation

For practical systems:

- d_0 : Single-query leakage baseline (empirically measured)
- a_S, a_M : Channel-specific leakage rates (from profiling)
- $d_{\max}$: Threat model bound (e.g., 10^6 queries/day)

5 Design Principles and Budget Management

5.1 The Allocation Problem

Given total budget $C_T = C_S + C_M < H(X)$, how should we allocate between privacy protection and delegation capacity?

5.2 Practical Budget Estimation

Monte Carlo Estimation:

Listing 1: Budget estimation

```
1 # Budget estimation
2 def estimate_mutual_info(samples_X, samples_Y):
3     # Use MINE or InfoNCE estimators
4     # Return confidence interval
5     return I_estimate, confidence_bounds
```

Runtime Monitoring:

- Track cumulative information release
- Implement privacy ledger similar to differential privacy
- Trigger alerts when approaching budget limits

5.3 Adaptive Control Framework

Listing 2: Adaptive budget controller

```
1 # Adaptive budget controller
2 class AdaptiveBudgetController:
3     def __init__(self, C_S_max, C_M_max):
4         self.budget_S = C_S_max
5         self.budget_M = C_M_max
6
7     def adjust_granularity(self, current_usage):
8         if current_usage > 0.8 * self.budget_S:
9             # Reduce disclosure precision
10            return reduced_precision_mode
```

5.4 Deployment Considerations

Handling Feedback Loops:

- Implement causal isolation barriers
- Use temporal batching to break correlations
- Apply differential privacy within S 's operations when needed

Active Adversary Mitigations:

- Rate limiting on queries
- Anomaly detection for unusual patterns
- Cryptographic commitments to prevent manipulation

6 Architectural Implementation

6.1 Measurement Design Principles

For Privacy Agent (S):

- Differential privacy within budget C_S
- Verifiable predicates using zero-knowledge proofs
- k -anonymity for categorical data
- Secure sketches for approximate matching

6.1.1 ZKP-Enhanced Selective Disclosure

The Swordsman's selective disclosure can be implemented using modern zero-knowledge proof systems:

- **Groth16:** For predicate verification with $O(1)$ proof size
 - Constant-size proofs (3 group elements)
 - Fast verification ($\sim 2\text{ms}$)

- Requires trusted setup per circuit
- **PLONK:** For circuit-based privacy constraints
 - Universal trusted setup
 - Flexible constraint systems
 - Larger proofs but better composability
- **Nova:** For recursive composition of privacy checks
 - No trusted setup
 - Efficient recursive proof composition
 - Enables cumulative budget verification

Concrete Construction: For attribute disclosure, the Swordsman generates:

$$\pi_{\text{attr}} = \text{ZKP}\{y_S = E_S(x) \wedge f(y_S) = 1\}$$


where f is a public predicate (e.g., “age ≥ 18 ”) and the proof reveals nothing beyond the predicate’s truth value.

For Delegation Agent (M):

- Minimize observation precision
- Batch queries to amortize leakage
- Caching to avoid redundant observations
- Query rate limiting

6.2 Separation Enforcement

Hardware-Based:



Software-Based:

- Container isolation with syscall filtering
- Formal verification of information flows
- Capability-based access control
- Audit logging for violations

6.2.1 Cryptographic Separation Verification

Rather than relying solely on trusted isolation, we can prove separation cryptographically:

$$\pi_{\text{sep}} = \text{ZKP}\{(Y_S \perp Y_M) \mid X\}$$

Using techniques from zero-knowledge proof systems:

- Commit to both observations: $c_S = \text{Commit}(Y_S, r_S)$, $c_M = \text{Commit}(Y_M, r_M)$
- Prove conditional independence via joint distribution commitments
- Enable third-party verification of architectural compliance

This approach enables:

- **Auditability:** External parties can verify separation without observing raw data
- **Composability:** Proofs can be recursively composed for multi-agent systems
- **Trustless Enforcement:** No need for trusted hardware or monitors

6.3 Isolation Verification

Formal Verification Approaches:

- Non-interference proofs using tools like FlowDroid
- Information flow security verification
- Covert channel capacity analysis

Runtime Instrumentation:

- Monitor timing patterns for correlation
- Detect shared resource usage
- Track API call sequences

6.3.1 ZKP-Based Budget Compliance

Zero-knowledge proofs enable verifiable budget tracking without revealing actual observations:

Protocol:

- Agent commits to observation: $c_i = \text{Commit}(y_i, r_i)$
- Proves cumulative budget compliance:

$$\pi_{\text{budget}} = \text{ZKP} \left\{ \sum I(X; y_i) \leq C_S \right\}$$

- Verifier checks without learning $\{y_i\}$

This addresses the monitoring challenge with cryptographic rather than trusted enforcement.

6.4 Zero-Knowledge Implementation Patterns

6.4.1 Range Proofs for Continuous Variables

For continuous variables, prove $y_S \in [a, b]$ without revealing exact value using Bulletproofs or similar range proof systems.

Example: Location disclosure

$$\pi_{\text{loc}} = \text{ZKP}\{\text{distance}(y_S, \text{target}) < 5\text{km}\}$$

6.4.2 Set Membership Proofs

Prove attribute membership in approved sets without revealing which element using Merkle tree commitments and Groth16.

Example: Credential verification

$$\pi_{\text{cred}} = \text{ZKP}\{y_S \in \text{AccreditedUniversities}\}$$

6.4.3 Predicate Verification

Prove arbitrary predicates $f(Y_S) = 1$ using circuit-based SNARKs (PLONK).

Example: Age verification

$$\pi_{\text{age}} = \text{ZKP}\{\text{age}(y_S) \geq 18 \wedge \text{age}(y_S) < 120\}$$

6.4.4 Cumulative Budget Tracking

Maintain cryptographic ledger:

$$\text{Ledger} = \{(c_i, \pi_i)\} \text{ for } i = 1 \text{ to } t$$

where c_i commits to observation i and π_i proves incremental budget consumption stays below C_S .

Using Nova for recursive composition:

$$\pi_1 = \text{ZKP}\{I(X; y_1) \leq C_S\} \tag{10}$$

$$\pi_t = \text{ZKP}\{\pi_{t-1} \wedge I(X; y_1, \dots, y_t) \leq C_S\} \tag{11}$$

6.5 Implementation Checklist

Pre-deployment:

- ☐ Estimate $H(X)$ for target domain
- ☐ Set $C_S + C_M \leq 0.7 \cdot H(X)$ (safety margin)
- ☐ Implement separation enforcement
- ☐ Verify isolation properties
- ☐ Deploy monitoring infrastructure

Runtime:

- ☐ Track actual $I(X; Y_S)$ and $I(X; Y_M)$
- ☐ Monitor separation violations
- ☐ Log reconstruction attempts
- ☐ Adjust budgets adaptively

7 Theoretical Predictions

7.1 Golden Ratio Convergence Hypothesis

Theoretical Prediction: Under certain optimization conditions, dual-agent systems may converge to allocation ratios near the golden ratio $\phi \approx 1.618$.

Theoretical Foundation: Consider value functions of the form $V(C_S, C_M) = U(C_M) \cdot P(C_S) \cdot \Psi(R_{\max})$ where:

- $U(C_M)$: Delegation utility (e.g., power-law: $U(x) = x^\alpha$)
- $P(C_S)$: Privacy protection value (e.g., $P(x) = x^\beta$)
- $\Psi(R_{\max})$: Reconstruction penalty

For Lagrangian optimization with power-law utilities, the optimal ratio satisfies:

$$\frac{\partial V / \partial C_S}{\partial V / \partial C_M} = \frac{\beta}{\alpha}$$

When $\beta/\alpha \approx \phi$ (protection slightly exceeds delegation), systems may show enhanced stability. This theoretical prediction suggests ϕ as a potential optimization attractor under symmetric importance of privacy and utility.

Status: This is a theoretical prediction derived from optimization analysis. Rigorous mathematical proof of universal convergence remains an open problem requiring further validation.

7.2 Open Research Question: Data Collection Needed

Research Opportunity: We seek to collect empirical data on C_S/C_M ratio patterns across different implementation domains to test the golden ratio convergence hypothesis. This is purely theoretical at present—no implementations or observations exist yet.

Call for Collaboration: Researchers and developers implementing dual-agent privacy systems are invited to share their allocation patterns, system behaviors, and empirical measurements. This data would be invaluable for validating or refuting the theoretical predictions presented in this paper.

7.3 Tetrahedral Emergence Hypothesis

Theoretical Prediction: Sustained S - M separation may naturally generate two additional measurement properties, forming a tetrahedral system architecture:

- **Reflect (R):** Temporal accumulation of S 's boundary decisions, creating memory properties
- **Connect (C):** Network effects from M 's delegation patterns, creating relational properties

Mathematical Formulation: If $(Y_S \perp Y_M) \mid X$ is maintained over time, the system may generate:

$$Y_R = R(Y_S^{1:t}, \tau) \quad [\text{Memory from S history}] \quad (12)$$

$$Y_C = C(Y_M^{1:t}, G) \quad [\text{Network from M interactions}] \quad (13)$$

By data processing inequality, $I(X; Y_R) \leq I(X; Y_S)$ and $I(X; Y_C) \leq I(X; Y_M)$, suggesting bounded additional leakage if these properties emerge.

Theoretical Rationale:

- **Minimum for sovereignty:** Two agents create separation, two properties enable persistence
- **Natural dimensions:** Temporal (R) and spatial (C) emerge from operational (S, M)
- **Information conservation:** Additional properties would be bounded by primary agent budgets

Status: This is a theoretical prediction about potential system evolution. Whether such emergence is inevitable, optimal, or even desirable requires further mathematical analysis and empirical validation.

7.4 Testable Predictions

If tetrahedral emergence occurs in real systems, we would expect to observe:

- **Temporal patterns:** Systems may develop memory/logging behaviors (potential “Reflect” property)
- **Network effects:** Inter-agent communication patterns may emerge (potential “Connect” property)

Important: These are theoretical predictions, not observations. No implementations exist yet to test whether such emergence actually occurs. These predictions provide concrete hypotheses for future experimental work.

7.5 Future Research Directions

- **Statistical Validation:**
 - Large-scale experiments across domains
 - Hypothesis testing for allocation patterns
 - Control for confounding variables
- **Theoretical Investigation:**
 - Derive optimal allocation from first principles
 - Characterize conditions for specific ratio convergence
 - Prove or disprove emergence hypotheses
- **Practical Extensions:**
 - Multi-agent generalizations (beyond two)
 - Dynamic budget adjustment algorithms
 - Integration with existing privacy tools

8 Theoretical Predictions and Open Questions

STATUS: PURELY THEORETICAL — This section presents unproven mathematical conjectures and theoretical predictions. No implementations, empirical data, or observations exist. These are research hypotheses requiring both formal mathematical proof and experimental validation.

8.1 Golden Ratio Hypothesis (Unproven)

Conjecture 8.1 (Golden Ratio Optimality - UNPROVEN). There may exist an optimization principle that drives optimal allocation ratios toward $\phi \approx 1.618$.

Theoretical Motivation:

If a joint optimization objective exists of the form:

$$\max_{C_S, C_M} U(C_M) \cdot P(C_S, C_M) \quad \text{s.t.} \quad C_S + C_M = B$$

where U is utility and P is privacy protection, then under certain smoothness and monotonicity conditions, the optimal ratio might be $C_S/C_M = \phi$.

What Would Be Required to Validate This:

- Formal mathematical proof deriving ϕ from first principles
- Precise definitions of utility and privacy functions
- Proof of existence and uniqueness of optimal ratio
- Characterization of exact conditions under which ϕ emerges
- Experimental implementations across multiple domains
- Statistical validation with large sample sizes

Current Status: Pure conjecture. No proof exists. No data exists. This is a mathematical hypothesis awaiting investigation.

8.2 Budget Allocation: Testable Framework

Proposed Metrics (for future experimental work):

Define efficiency metrics for implementations:

- **Privacy Efficiency:** $\eta_P = \text{Privacy Achieved} / \text{Budget Spent}$
- **Utility Efficiency:** $\eta_U = \text{Utility Achieved} / \text{Budget Spent}$
- **Joint Efficiency:** $\eta_J = \eta_P \cdot \eta_U$

Theoretical Predictions (to be tested):

- η_P may increase with C_S but with diminishing returns
- η_U may increase with C_M but with diminishing returns
- η_J might show local maxima near $C_S/C_M \approx \phi$ if the golden ratio hypothesis holds

Critical Issues Requiring Resolution:

- “Privacy Achieved” and “Utility Achieved” need rigorous, measurable definitions
- Metrics may be highly implementation-dependent
- Optimization landscape likely varies by application domain
- No implementations exist yet to test these predictions

9 Tetrahedral Emergence Hypothesis

STATUS: HIGHLY SPECULATIVE — This section presents a geometric interpretation that has not been mathematically validated.

9.1 Four-Force Model

Conjecture 9.1 (Tetrahedral Structure). The dual-agent system, when operating over time, may exhibit four distinct information flows:

- **Swordsman (S)**: Privacy enforcement (boundary-making)
- **Mage (M)**: Delegation (capability projection)
- **Reflect (R)**: Temporal integration (memory formation)
- **Connect (C)**: Network coordination (relationship building)

Geometric Interpretation:

If these four forces exist and maintain pairwise relationships, they may form a tetrahedral structure where:

- Each vertex represents one force
- Each edge represents an information channel
- The structure is minimal (no redundant channels)
- Stability requires specific geometric relationships

Mathematical Formulation (Tentative):

Let I_{ij} denote mutual information between forces i and j . The tetrahedron hypothesis suggests:

$$I_{SM} = 0 \quad (\text{proven separation}) \tag{14}$$

$$I_{SR} \approx \phi \cdot I_{MR} \quad (\text{conjectured}) \tag{15}$$

$$I_{SC} \approx \phi \cdot I_{MC} \quad (\text{conjectured}) \tag{16}$$

$$I_{RC} \approx \phi^2 \quad (\text{highly speculative}) \tag{17}$$

9.2 Emergence from Dual-Agent Core

Claim 8.1 (UNPROVEN): Reflect and Connect are not separate agents but emergent properties of sustained S - M separation over time.

Proposed Mechanism:

- **Reflect emerges** when Swordsman maintains consistent boundaries across sessions, creating temporal correlation structure
- **Connect emerges** when Mage builds relationships with other agents, creating network correlation structure

Required Validation:

- Formal definition of “emergence” in this context
- Proof that R and C are inevitable given S - M separation
- Demonstration that R and C cannot be reduced to S and M
- Empirical observation in deployed systems

9.3 Golden Ratio in Tetrahedral Context

Speculation 8.1: If the four-force tetrahedron exists and maintains stability, the golden ratio may govern relationships:

$$I_S/I_M = \phi, \quad I_S/I_R = \phi^2, \quad I_S/I_C = \phi^3$$

Geometric Justification (Informal):

The golden ratio appears in regular geometric structures (pentagons, Fibonacci spirals). If information flows in the tetrahedron seek geometric efficiency, ϕ -based relationships might provide optimal packing or stability properties.

Critical Acknowledgment: This geometric interpretation is highly speculative and serves primarily as a framework for future investigation. No rigorous derivation exists, and the hypothesis may ultimately prove incorrect.

10 Discussion

10.1 Summary of Proven Guarantees

We have rigorously established:

- Separation enables additive mutual information bounds
- Combined with budgets, guarantees $R_{\max} < 1$
- Fano’s inequality ensures minimum error rates
- Approximate separation degrades gracefully
- ZKP constructions enable cryptographic enforcement

These results hold unconditionally, independent of computational assumptions (except for ZKP implementations which require standard cryptographic hardness).

10.2 Practical Implementation Pathways

The framework can be deployed through:

- **Hybrid Approaches:** Combine with differential privacy for additional protection
- **Layered Architecture:** Use TEEs for hardware isolation, containers for software
- **Gradual Adoption:** Start with low-sensitivity applications, expand as confidence grows

10.3 Relationship to Existing Privacy Frameworks

Framework	Focus	Our Approach	Synergy
Differential Privacy	Statistical noise	Structural separation	Use DP within S
Secure MPC	Distributed computation	Distributed observation	Complementary
Information Flow Control	Taint tracking	Quantitative bounds	Enhanced metrics

Table 1: Comparison with existing privacy frameworks

10.4 Open Questions and Theoretical Predictions

Unresolved Questions:

- Is two-agent separation optimal for privacy-utility trade-offs?
- Can the golden ratio convergence be proven from first principles?
- Does tetrahedral emergence occur naturally or require explicit design?
- How do these bounds compose in multi-agent systems?

Theoretical Predictions Requiring Validation:

- Golden ratio proportions as optimization attractors in allocation strategies
- Tetrahedral structure as natural consequence of sustained dual-agent separation
- Emergent properties R and C as inevitable system evolution
- Four-force structure as minimal complete sovereignty architecture

10.5 Limitations and Assumptions

Key Limitations:

- **Conditional Independence:** Hard to enforce perfectly in practice
- **Passive Adversary:** Active attacks not fully addressed
- **Known Distributions:** Uncertainty in $P(X)$ affects budgets
- **Static Budgets:** Dynamic environments may require adaptation

Critical Assumptions:

- Architectural isolation is maintainable
- Side channels are bounded and detectable
- Budget estimation is sufficiently accurate

10.6 Experimental Roadmap

Immediate Next Steps:

- Implement reference architecture with common privacy benchmarks
- Develop standardized test suite for separation verification
- Create open-source monitoring tools

Medium-term Goals:

- Deploy in real applications with user studies
- Validate across diverse domains
- Establish best practices for budget setting

Long-term Research:

- Prove or refute optimal allocation theorems
- Extend to multi-agent and distributed settings
- Integrate with emerging privacy technologies

11 Related Extended Work

11.1 Connections to Complex Systems

The potential emergence of structure from simple rules connects to:

- **Self-organization** [17]
- **Emergent complexity** [18]
- **Scale-free networks** [19]
- **Allometric scaling** [20]

11.2 Privacy Technology Integration

Our framework complements:

- **Zero-knowledge proofs:** Can implement S 's selective disclosure with cryptographic guarantees (Groth16, PLONK, Nova)
- **Secure enclaves:** Hardware enforcement of separation
- **Homomorphic encryption:** Computation within M 's bounds
- **Privacy pools:** Network effects without individual exposure

11.3 Economic and Game Theory

If golden ratios emerge naturally:

- **Nash equilibria** might converge to ϕ -proportions
- **Market mechanisms** could discover optimal allocations
- **Evolutionary pressure** might select for tetrahedral structures

Economic Enforcement of Separation:

The architectural separation $(Y_S \perp Y_M) \mid X$ can be economically enforced through dual-token markets:

- SWORD tokens earned exclusively through Swordsman chronicles (privacy domain)
- MAGE tokens earned exclusively through Mage chronicles (delegation domain)
- Market separation creates economic pressure against agent merger
- Budget tracking C_S and C_M verifiable on-chain
- Guardian staking (10,000 SWORD) maintains collective compression standards

If a First Person attempts to merge agents (violating $(Y_S \perp Y_M) \mid X$), they lose earning capability in both domains—creating self-enforcing separation through incentive alignment.

Signal-Based Sustainability:

Rather than speculative token sales, sustainable funding emerges from ceremony and signal fees:

- Genesis ceremony: 1 ZEC (\$500 at \$500/ZEC) creates agent pair once per ecosystem

- Ongoing signals: 0.01 ZEC (\$5) each, continuous proof-of-comprehension
- Fee distribution: 61.8% transparent pool, 38.2% shielded pool (internal allocation per ecosystem)
- Self-sustaining through activity-based revenue

Relationship to Proven Results:

The economic model described above is *one possible implementation*—the mathematical guarantees established in Sections 3–6 hold independent of tokenomic choices. The reconstruction ceiling, error floor, and separation bounds are proven theorems that remain valid regardless of:

- Token model (dual, single, or no tokens)
- Funding mechanism (ceremonies, grants, or other)
- Reward structure (progressive, fixed, or algorithmic)
- Network effects (VRCs optional, not required)

What the proven mathematics *enables* is sustainable tokenomics grounded in information theory rather than speculation. The budget constraint $C_S + C_M < H(X)$ provides a scarcity bound; how that bound is implemented economically is a design choice.

For complete economic architecture, see companion document: “VRC Protocol: Economic Architecture”

12 Conclusion

We have established rigorous information-theoretic bounds for dual-agent privacy architectures with enforced separation. The proven results—additive mutual information under separation, reconstruction ceilings below unity, and guaranteed error floors—provide solid foundations for privacy-preserving agent systems.

We integrate zero-knowledge proof systems as core implementation primitives, providing concrete constructions using Groth16, PLONK, and Nova protocols. This enables cryptographic rather than merely architectural enforcement of separation and budget constraints, offering stronger guarantees in adversarial settings.

The key insight remains powerful: structural separation with budget constraints creates fundamental privacy guarantees independent of computational assumptions. This offers a principled approach to the agent privacy problem that complements existing methods.

We present theoretical conjectures about golden ratio optimization and tetrahedral emergence, but emphasize these remain unproven mathematical hypotheses with no empirical validation. The framework’s value lies entirely in its proven information-theoretic guarantees and the implementation roadmap it provides for future builders.

Version Statement

Version 3.2: This edition clearly separates proven information-theoretic results from unproven theoretical conjectures. The core separation bounds, reconstruction ceilings, and error guarantees are rigorously established through formal proof. Golden ratio hypotheses and tetrahedral emergence remain purely theoretical conjectures with no empirical validation—no implementations or observations exist yet. We explicitly distinguish mathematical theorems from speculative predictions and invite collaboration on both formal proof development and experimental validation. For implementation guidance, see the companion *Oxagentprivacy Whitepaper*.

Acknowledgments

To the 0xagentprivacy project for exploring privacy-first infrastructure. To BGIN for discussions on AI governance. To Kwaai for advancing personal AI sovereignty. To First Person Project for pioneering user-centric identity frameworks. To the reviewers whose feedback strengthened this work.

References

- [1] Dwork, C. & Roth, A. (2014). The Algorithmic Foundations of Differential Privacy. *Foundations and Trends in Theoretical Computer Science*.
- [2] Goldreich, O. (2004). Foundations of Cryptography: Volume 2, Basic Applications. Cambridge University Press.
- [3] Sabelfeld, A. & Myers, A. C. (2003). Language-based information-flow security. *IEEE Journal on Selected Areas in Communications*, 21(1), 5–19.
- [4] Kairouz, P., Oh, S., & Viswanath, P. (2017). The Composition Theorem for Differential Privacy. *IEEE Transactions on Information Theory*, 63(6), 4037–4039.
- [5] Shannon, C. E. (1948). A Mathematical Theory of Communication. *Bell System Technical Journal*.
- [6] Cover, T. & Thomas, J. (2006). Elements of Information Theory. Wiley.
- [7] Fano, R. M. (1961). Transmission of Information. MIT Press.
- [8] Millen, J. K. (1987). Covert Channel Capacity. *IEEE Symposium on Security and Privacy*.
- [9] Gray III, J. W. (1993). On Analyzing the Bus-Contention Channel. *IEEE Computer Security Foundations Workshop*.
- [10] Giles, J. & Hajek, B. (2002). An Information-Theoretic and Game-Theoretic Study of Timing Channels. *IEEE Transactions on Information Theory*, 48(9), 2455–2477.
- [11] Kocher, P., Jaffe, J., & Jun, B. (1999). Differential Power Analysis. *CRYPTO '99*.
- [12] Yarom, Y. & Falkner, K. (2014). FLUSH+RELOAD: A High Resolution, Low Noise, L3 Cache Side-Channel Attack. *USENIX Security*.
- [13] Mangard, S., Oswald, E., & Popp, T. (2007). Power Analysis Attacks: Revealing the Secrets of Smart Cards. Springer.
- [14] Murdoch, S. J. & Danezis, G. (2005). Low-Cost Traffic Analysis of Tor. *IEEE Symposium on Security and Privacy*.
- [15] Dwork, C., Rothblum, G. N., & Vadhan, S. (2010). Boosting and Differential Privacy. *FOCS 2010*.
- [16] Mironov, I. (2017). Rényi Differential Privacy. *IEEE Computer Security Foundations Symposium*.
- [17] Prigogine, I. & Stengers, I. (1984). Order Out of Chaos. Bantam Books.
- [18] Kauffman, S. (1993). The Origins of Order: Self-Organization and Selection in Evolution. Oxford University Press.

- [19] Barabási, A. L. (2002). *Linked: The New Science of Networks*. Perseus Books.
- [20] West, G. (2017). *Scale: The Universal Laws of Life and Death in Organisms, Cities and Companies*. Penguin.
- [21] Ball, P. (2009). *Nature’s Patterns: A Tapestry in Three Parts*. Oxford University Press.
- [22] Livio, M. (2002). *The Golden Ratio: The Story of Phi*. Broadway Books.
- [23] Groth, J. (2016). On the Size of Pairing-based Non-interactive Arguments. *EUROCRYPT 2016*. ePrint Archive 2016/260.
- [24] Gabizon, A., Williamson, Z. J., & Ciobotaru, O. (2019). PLONK: Permutations over Lagrange-bases for Oecumenical Noninteractive arguments of Knowledge. ePrint Archive 2019/953.
- [25] Kothapalli, A., Setty, S., & Tzialis, I. (2021). Nova: Recursive Zero-Knowledge Arguments from Folding Schemes. ePrint Archive 2021/370.

A Proof Details

A.1 Complete Chain Rule Expansion

For completeness, the full chain rule for four variables:

$$I(X; Y_S, Y_M, Y_R, Y_C) = I(X; Y_S) + I(X; Y_M | Y_S) + I(X; Y_R | Y_S, Y_M) + I(X; Y_C | Y_S, Y_M, Y_R)$$

Each conditional term is bounded by the unconditional under independence assumptions.

A.2 Fano’s Inequality Tightness

For binary classification with symmetric channels, Fano’s bound is nearly tight. For larger alphabets, the bound loosens but remains non-trivial for practical entropy values.

B Implementation Pseudocode

B.1 Basic Dual-Agent System

Listing 3: Basic dual-agent system

```

1 # Basic dual-agent system
2 class DualAgentPrivacy:
3     def __init__(self, entropy_bits, safety_factor=0.7):
4         self.H_X = entropy_bits
5         self.budget = self.H_X * safety_factor
6
7         # Allocation (example values, not empirically validated)
8         # Golden ratio hypothesis suggests C_S ~= 0.62, C_M ~= 0.38
9         self.C_S = self.budget * 0.62
10        self.C_M = self.budget * 0.38
11
12    def enforce_separation(self):
13        """Implementation-specific isolation"""
14        # Use TEEs, containers, or formal verification
15        pass
16
17    def measure_leakage(self):

```

```

18         I_S = self.estimate_mutual_info(self.Y_S, self.X)
19         I_M = self.estimate_mutual_info(self.Y_M, self.X)
20         I_joint = self.estimate_mutual_info((self.Y_S, self.Y_M), self.
21             X)
22
23         # Verify separation
24         separation_violation = I_joint - I_S - I_M
25         return I_S, I_M, separation_violation

```

B.2 Adaptive Budget Controller

Listing 4: Adaptive budget controller

```

1 # Adaptive budget controller
2 class AdaptiveBudgetController:
3     def __init__(self, total_budget, initial_ratio=1.5):
4         self.total = total_budget
5         self.ratio = initial_ratio
6         self.history = []
7
8     def update(self, utility_score, privacy_score):
9         """Adjust allocation based on performance"""
10        self.history.append((utility_score, privacy_score))
11
12        if len(self.history) > 100:
13            # Analyze trends, adjust ratio
14            if utility_score < threshold:
15                self.ratio *= 0.9 # Give more to Mage
16            elif privacy_score < threshold:
17                self.ratio *= 1.1 # Give more to Swordsman
18
19    def get_budgets(self):
20        C_S = self.total * (self.ratio / (1 + self.ratio))
21        C_M = self.total * (1 / (1 + self.ratio))
22        return C_S, C_M

```

B.3 Testing Separation Violations

Listing 5: Separation violation testing

```

1 # Separation violation testing
2 def test_separation_violation(system, num_samples=10000):
3     """Empirically verify conditional independence"""
4     samples = []
5     for _ in range(num_samples):
6         x = sample_secret()
7         y_s, y_m = system.observe(x)
8         samples.append((x, y_s, y_m))
9
10    # Compute  $I(Y_S; Y_M | X)$ 
11    violation = estimate_conditional_mi(samples)
12
13    # Check against threshold
14    if violation > epsilon:
15        return "VIOLATION DETECTED", violation
16    return "SEPARATION MAINTAINED", violation

```

B.4 Budget Compliance Monitoring

Listing 6: Budget compliance monitoring

```
1 # Budget compliance monitoring
2 def monitor_budget_compliance(agent, budget, window=1000):
3     """Real-time budget monitoring"""
4     cumulative_info = 0
5     measurements = deque(maxlen=window)
6
7     while True:
8         x, y = agent.get_next_observation()
9         instant_mi = estimate_instant_mi(x, y)
10        measurements.append(instant_mi)
11
12        cumulative_info = sum(measurements)
13        if cumulative_info > budget:
14            agent.trigger_throttling()
15            log_budget_violation()
```